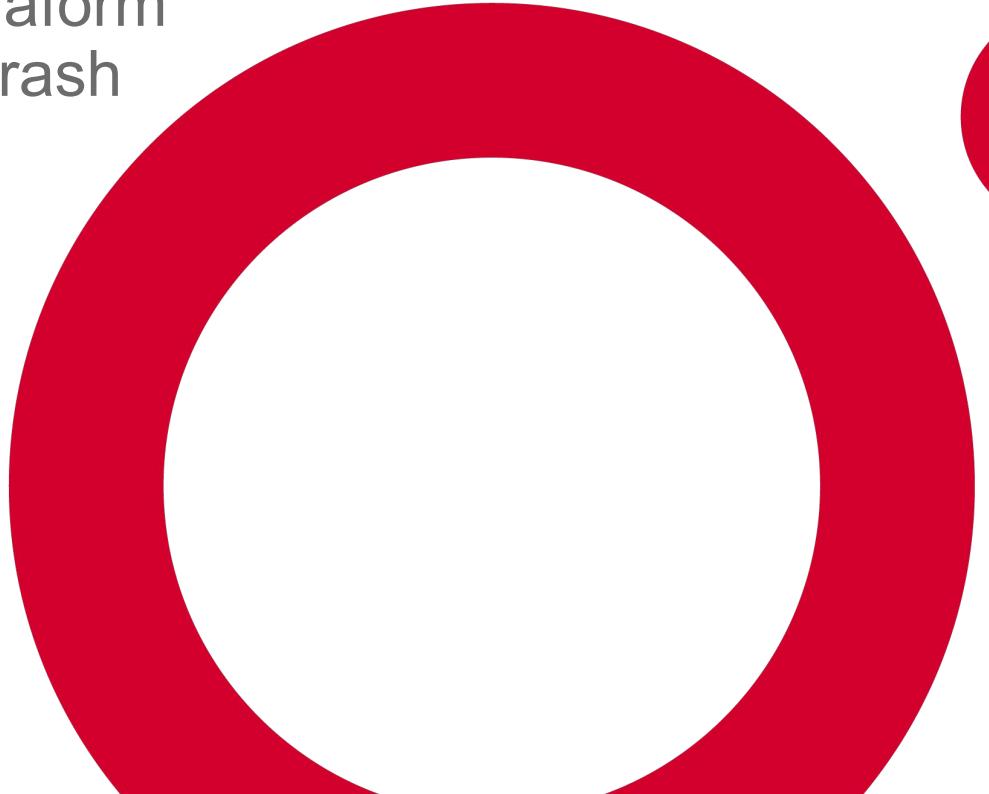




HashiCorp Certified: Terraform Associate Certification Crash Course

Curriculum Version 004



About the trainer



bmuschko@automatedascent.com



benjaminmuschko



bmuschko



automatedascent.com

Certification Exam

Objectives, Curriculum, Prerequisites, Learning Resources

Exam Objectives

“Basic understanding of concepts and skills associated with open source HashiCorp Terraform and HCP Terraform.”



The certification program allows users to demonstrate their competence in a multiple choice test.

<https://www.hashicorp.com/certification/terraform-associate>



Curriculum Changes

Version 004 replaced 003 in January 2026

Content Differences Between the (003) and (004) Exams

Exam updates summary:

- Four new topics
- Tests on Terraform version 1.12
- Includes HCP Terraform content

New topics covered in (004):

- 4f: `depends_on` and `create_before_destroy` lifecycle rules
- 4g: Validate configuration using custom conditions
- 4h: Ephemeral values and write-only arguments
- 8c: Describe how to organize and use HCP Terraform workspaces and projects

<https://www.hashicorp.com/certification/terraform-associate-004>



The Curriculum

1	Infrastructure as Code (IaC) with Terraform
1a	Explain what IaC is
1b	Describe the advantages of IaC patterns
1c	Explain how Terraform manages multi-cloud, hybrid cloud, and service-agnostic workflows
2	Terraform fundamentals
2a	Install and version Terraform providers
2b	Describe how Terraform uses providers
2c	Write Terraform configuration using multiple providers
2d	Explain how Terraform uses and manages state
3	Core Terraform workflow
3a	Describe the Terraform workflow
3b	Initialize a Terraform working directory
3c	Validate a Terraform configuration
3d	Generate and review an execution plan for Terraform
3e	Apply changes to infrastructure with Terraform
3f	Destroy Terraform-managed infrastructure
3g	Apply formatting and style adjustments to a configuration



The Curriculum

3	Core Terraform workflow
3a	Describe the Terraform workflow
3b	Initialize a Terraform working directory
3c	Validate a Terraform configuration
3d	Generate and review an execution plan for Terraform
3e	Apply changes to infrastructure with Terraform
3f	Destroy Terraform-managed infrastructure
3g	Apply formatting and style adjustments to a configuration
4	Terraform configuration
4a	Use and differentiate <code>resource</code> and <code>data</code> blocks
4b	Refer to resource attributes and create cross-resource references
4c	Use variables and outputs
4d	Understand and use complex types
4e	Write dynamic configuration using expressions and functions
4f	Define resource dependencies in configuration
4g	Validate configuration using custom conditions
4h	Understand best practices for managing sensitive data, including secrets management with Vault



The Curriculum

5	Terraform modules
5a	Explain how Terraform sources modules
5b	Describe variable scope within modules
5c	Use modules in configuration
5d	Manage module versions
6	Terraform state management
6a	Describe the local backend
6b	Describe state locking
6c	Configure remote state using the backend block
6d	Manage resource drift and Terraform state
7	Maintain infrastructure with Terraform
7a	Import existing infrastructure into your Terraform workspace
7b	Use the CLI to inspect state
7c	Describe when and how to use verbose logging
8	HCP Terraform
8a	Use HCP Terraform to create infrastructure
8b	Describe HCP Terraform collaboration and governance features
8c	Describe how to organize and use HCP Terraform workspaces and projects
8d	Configure and use HCP Terraform integration



Candidate Prerequisites



Basic Terminal Skills



Basic understanding of on premises and
cloud architecture



Preparing for the Exam

Documentation is not allowed during the test

- Learning Path:

[https://developer.hashicorp.com/terraform/tutorials/certification-004
/associate-study-004](https://developer.hashicorp.com/terraform/tutorials/certification-004/associate-study-004)

- Exam Content List:

[https://developer.hashicorp.com/terraform/tutorials/certification-004
/associate-review-004](https://developer.hashicorp.com/terraform/tutorials/certification-004/associate-review-004)



Preparing for the Exam

Documentation is not allowed during the test

- Sample Questions:

[https://developer.hashicorp.com/terraform/tutorials/certification-004
/associate-questions-004](https://developer.hashicorp.com/terraform/tutorials/certification-004/associate-questions-004)



Infrastructure as Code (IaC)

Concepts and Benefits

Infrastructure as Code (IaC)



Manage infrastructure with the help of code

- Treats all aspects of operations as software via configuration
- Code is tracked in a SCM repository
- Automation makes the provisioning process consistent, repeatable, and updates fast & reliable



Repeatable Process



Clear instructions that describe the desired state

- A set of instructions are defined with the help of a declarative language
- Operations are idempotent e.g. an update to the environment will only make necessary changes but not duplicate what already exists
- **Example:** A database was deployed with a port but the port needs to change. We can simply change the port number in the code we used to deploy the database in the first place.



Consistent Environments



Environment should look extremely similar

- Projects often use a variety of deployment environments e.g. development, staging, and production
- The same automation code can be used to provision infrastructure so that it looks consistent across all environments
- **Example:** Deploying an EC2 server across all environments but with slightly different regions



Reusable Functionality



Configuration can be abstracted and applied to a set of projects

- Configuration is defined with the help of code
- Code can be shared across different repositories
- **Example:** Deploying a HTTP web server for multiple projects with the same piece of code.



Self-Documenting



Source code represents the architecture

- Each piece of infrastructure has been described with a set of instructions
- No more guesswork on what configuration has been used to provision infrastructure
- **Example:** New engineering team members can read source code to understand how infrastructure is configured



Financial Savings



Increased efficiency, less mistakes through automation

- Reduced risk due to minimizing human error
- Infrastructure can be verified with automated tests
- IaC functions can be used to spin down environments during times of less traffic
- **Example:** Decrease the manual grunt work for DevOps personnel and spending it on mission-critical tasks instead



Q & A



Terraform 101

What is Terraform and How Does it Implement IaC?

What is Terraform?



Infrastructure automation tool

- Business source licensed (BSL) and cloud provider-agnostic
- Configuration is expressed in a declarative language with either HashiCorp Configuration Language (HCL) or JSON
- Deployment of infrastructure happens with a push-based approach (no agent to be installed on remote machines)



How Does Terraform Work?



Binary makes API calls to cloud providers

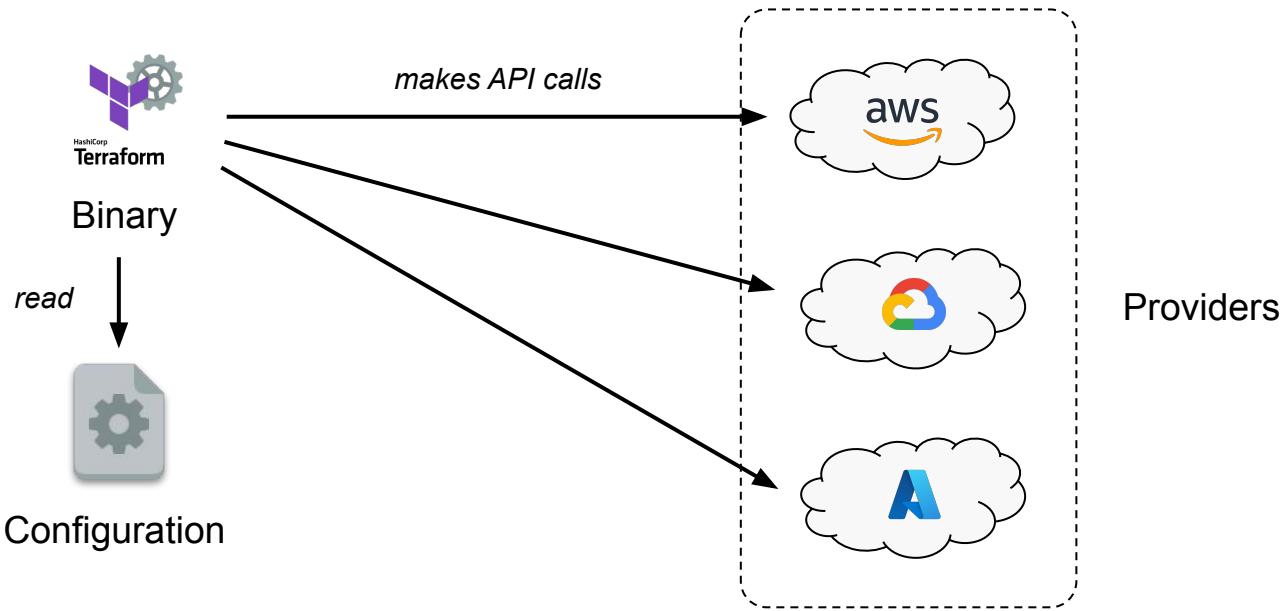
- CLI tool for deploying infrastructure to one or many cloud provider(s)
e.g. AWS, Azure, or Google Cloud
- Under the hood, makes API calls on behalf of a *provider* including authentication mechanisms
- Terraform *configurations* are the codified instructions in the form of a text file that tell it which API calls to make



Interaction with Cloud Providers



Binary makes API calls to cloud providers



Cloud Provider Portability



Features are different, Terraform's technical approach is not

- Cloud providers don't support the exact same infrastructure
- Terraform allows you to use the same approach to define provider-specific configuration
- You use the same Terraform language, toolset, and IaC practices



Terraform Components



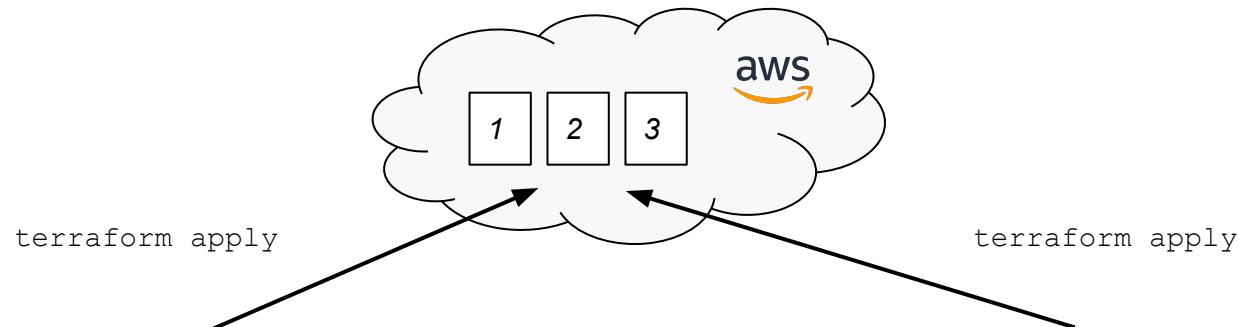
Key building blocks in architecture

- *Executable*: Binary run from the command line that contains Terraform's core functionality
- *Configuration file(s)*: Files with the extension `.tf` or `.tfvars` that define the desired configuration for provisioning infrastructure
- *Provider plugins*: Executables invoked by Terraform to interact with cloud provider APIs, hosted on a registry
- *State data*: The desired configuration and its current state



Persisting State

Knows what has been created before and applies changes



```
resource "aws_instance" "example" {  
  count      = 2  
  ami        = "ami-0c55b159cbfafef0"  
  instance_type = "t2.micro"  
}
```

```
resource "aws_instance" "example" {  
  count      = 3  
  ami        = "ami-0c55b159cbfafef0"  
  instance_type = "t2.micro"  
}
```



State Storage

Where and how is state stored and updated?

- Stores state in “internal database”, each resource is represented by an key-value pair in the entry
- Any changes to the resources will be reflected in the state
- Locally, the state is saved in the file `terraform.tfstate`
- Remote state handling exists to support consistent team collaboration



Benefits of State



Keeping state fulfills important requirements

- *Dependencies:* Resources can have dependencies on each other, Terraform retains this metadata to be able to safely perform operations e.g. delete
- *Performance:* Terraform stores a cache of the attribute values for all resources in the state for performance reasons
- *Consistency:* Terraform employs locking to avoid synchronization and collaboration issues



Terraform and OpenTofu



Differences and use cases

- In August 2023, HashiCorp changed Terraform's license from MPL to BSL. OpenTofu emerged as a community fork. The Terraform Associate certification is **not** about OpenTofu.
- Key differences emerging
 - OpenTofu added native state encryption
 - Feature development is diverging over time
 - Some newer Terraform features may not appear in OpenTofu (and vice versa)





Terraform and Ansible

Differences and use cases

- **Ansible:** Configuration management tool for installing/configuring software and tools in the infrastructure
- **Terraform:** May invoke Ansible after infrastructure has been deployed
- When to involve Ansible?
 - You already have Ansible playbooks you want to invoke
 - Nowadays, it's more common to create an immutable image that already contains the software needed



Terraform and Packer

Differences and use cases

- **Packer:** Creates machine images for multiple platforms
- **Terraform:** Uses machine image to provision infrastructure
- When to involve Packer?
 - To shorten deployment time by baking needed functionality into AMI
 - To simplify or avoid logic in Terraform needed for installing additional tooling, monitoring, logging etc. on top of base AMI





Terraform and Consul

Do I need to involve a state backend?

- **Terraform** can store state data in a backend to improve team collaboration
- **Consul** is one of the configurable backends supported by Terraform. Effectively, it's a key-value store installed on a dedicated server
- When to involve Consul?
 - When your team works on Terraform configuration in parallel and runs into potential conflicts a lot
 - If you do not want to buy into HCP Terraform



Q & A



Terraform Installation and Configuration

Binary Installation, AWS Account, IDE Support

Installing Terraform



Easy to install on all operating systems

- Manual installation
 - Download ZIP file containing the pre-compiled binary
 - Add binary to PATH environment variable
- Using a package manager
 - Available via Homebrew and Chocolatey
 - Takes care of adding binary to terminal



Verifying and Using Terraform

The `terraform` executable is the main entry point

```
$ terraform version
Terraform v1.14.1
on darwin_amd64
```

```
$ terraform -help
Usage: terraform [global options] <subcommand> [args]

The available commands for execution are listed below.
The primary workflow commands are given first, followed by
less common or more advanced commands.

Main commands:
  init          Prepare your working directory for other commands
  validate      Check whether the configuration is valid
  plan          Show changes required by the current configuration
  apply         Create or update infrastructure
  destroy       Destroy previously-created infrastructure

  ...


```



Switching Between Versions

tfswitch lets you manage and use versions in parallel



```
$ tfswitch 1.2.1
$ terraform version
Terraform v1.2.1
on darwin_amd64
```

```
$ tfswitch 1.3.2
$ terraform version
Terraform v1.3.2
on darwin_amd64
```





Setting Up an AWS Account

Create new account at <https://aws.amazon.com>

The screenshot shows the AWS homepage with a dark blue background. At the top right, there is a navigation bar with links for Contact Us, Support, English, My Account, and Sign In. A prominent orange button labeled "Create an AWS Account" is highlighted with a red box and an arrow pointing to it. Below the navigation bar, the text "Start Building on AWS Today" is displayed in white. A paragraph below it states: "Whether you're looking for compute power, database storage, content delivery, or other functionality, AWS has the services to help you build sophisticated applications with increased flexibility, scalability and reliability." A "Get Started for Free" button is located below this text. The page is divided into two main sections: "For Builders" and "For Decision Makers". The "For Builders" section includes icons for launching an application and connecting with others, along with descriptions for each. The "For Decision Makers" section includes icons for optimizing business value and reinventing with data, along with descriptions for each.





Setting Up an AWS Account

Free tier is sufficient for exercises in this course

aws

Sign up for AWS

Explore Free Tier products with a new AWS account.

To learn more, visit aws.amazon.com/free.

Root user email address
Used for account recovery and some administrative functions

AWS account name
Choose a name for your account. You can change this name in your account settings after you sign up.

Verify email address

OR

Sign in to an existing AWS account

Basic support - Free

- Recommended for new users just getting started with AWS
- 24x7 self-service access to AWS resources
- For account and billing issues only
- Access to Personal Health Dashboard & Trusted Advisor



Free Trials

Short-term free trial offers start from the date you activate a particular service



12 months free

Enjoy these offers for 12-months following your initial sign-up date to AWS



Always free

These free tier offers do not expire and are available to all AWS customers





Retrieving AWS Credentials

Create a new access key and store in safe place

The screenshot shows the AWS Management Console navigation bar. From left to right, it includes: a triangle icon, a user icon, 'N. Virginia ▾', 'bmuschko ▾', 'Account ID: 0314-0087-8386' with a copy icon, and a dark blue background area containing the following items:

- Account
- Organization
- Service Quotas
- Billing Dashboard
- Security credentials** (this item is highlighted with a red box)
- Settings

At the bottom right of this list is an orange 'Sign out' button.

The screenshot shows the 'Your Security Credentials' page. At the top, there's a header with three collapsed sections: 'Password', 'Multi-factor authentication (MFA)', and 'Access keys (access key ID and secret access key)' (the last one is expanded). Below the header, there's a note about using access keys for programmatic calls and a warning about sharing secret keys. A table at the bottom lists columns: 'Created', 'Access Key ID', 'Last Used', 'Last Used Region', 'Last Used Service', 'Status', and 'Actions'. At the bottom of the table is a blue 'Create New Access Key' button, which is also highlighted with a red box.





Setting Provider Credentials

In current shell or from typical credentials location

```
$ export AWS_ACCESS_KEY_ID=<access-key-id>
$ export AWS_SECRET_ACCESS_KEY=<secret-access-key>
```

or

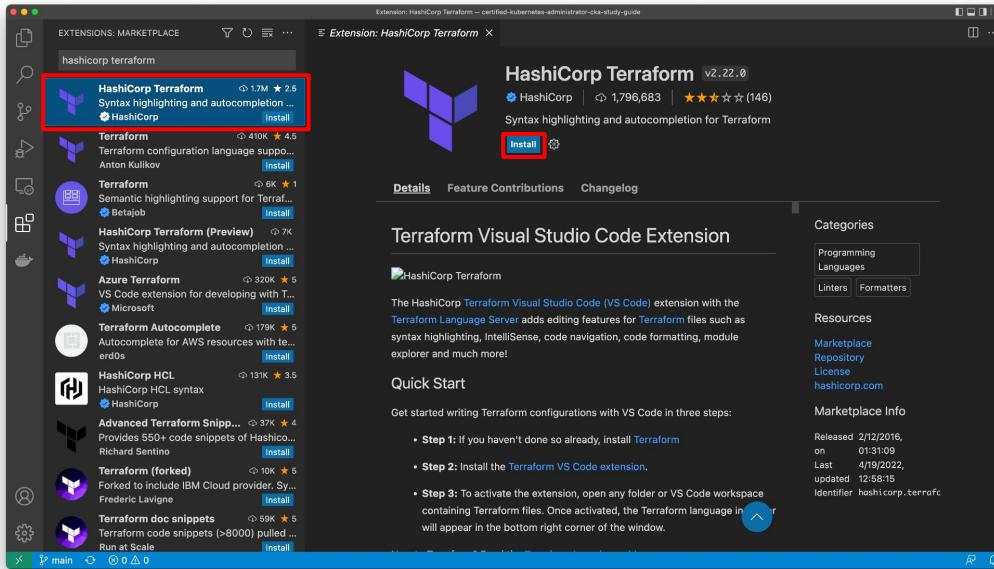
```
$ cat $HOME/.aws/credentials
[default]
aws_access_key_id=<access-key-id>
aws_secret_access_key=<secret-access-key>
```



IDE Integration with VSCode



Search for “HashiCorp Terraform” in the Extensions



EXERCISE

Installing Terraform
and Setting AWS
Credentials



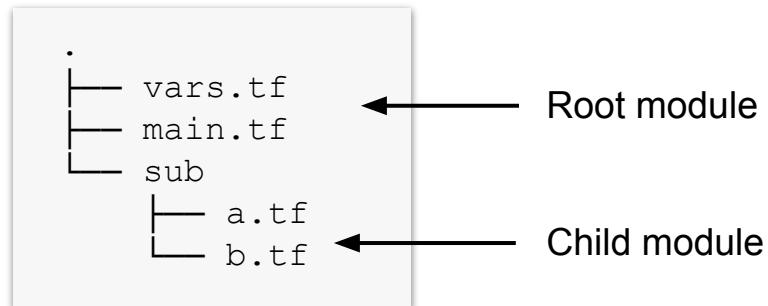
Getting Started With Terraform

Basic Concepts, Syntax, and Workflow

Files and Directories

Terraform configuration is defined by a collection of files

- *Configuration files* contain plain text instruction and have the extension `.tf` or `.tf.json`
- A *module* defines a set of (potentially versioned) configuration file(s)



Object Types

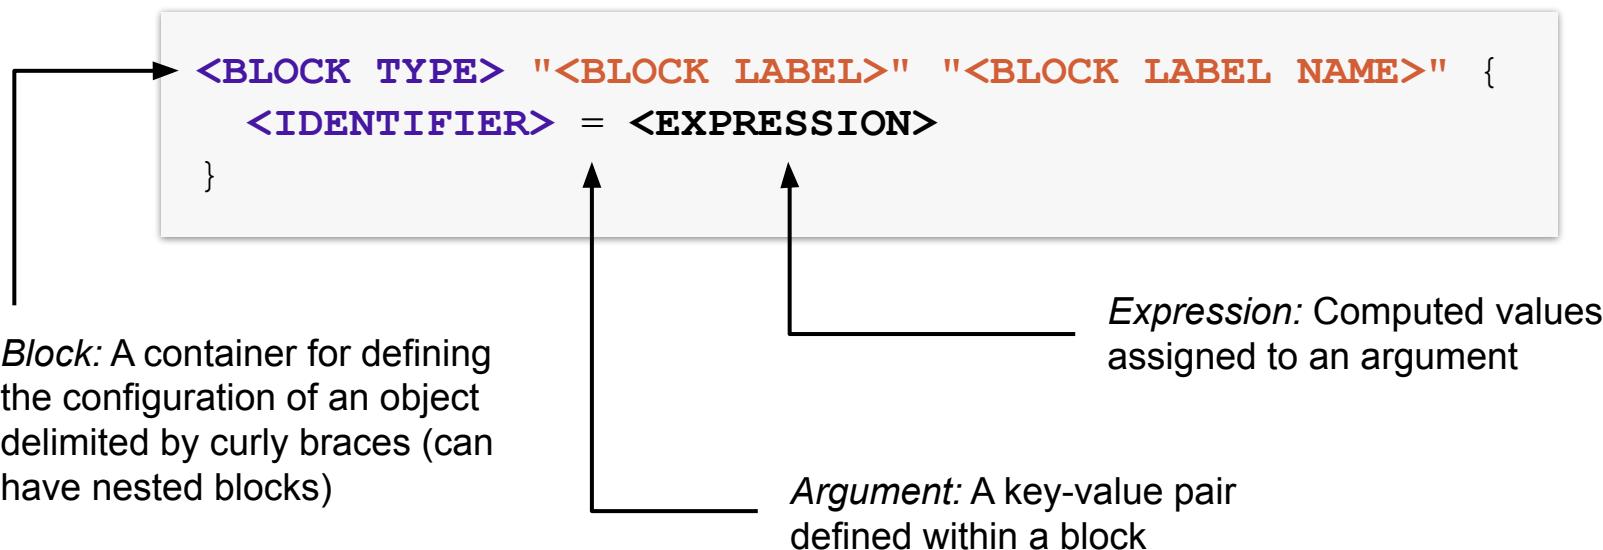
High-level elements used in a configuration file

- *Provider*: Allows Terraform to interact with a cloud provider through its API e.g. AWS or Azure, SaaS provider, or other APIs
- *Resource*: Defines the infrastructure pieces to be created in a target environment e.g. EC2 instance, a VPC, or a web server
- *Data source*: Can query information from a provider that can be used in the configuration e.g. a list of available AMIs



Configuration Syntax

HCL is preferred over JSON



Syntax Constraints

Just of a couple of little gotchas to look out for

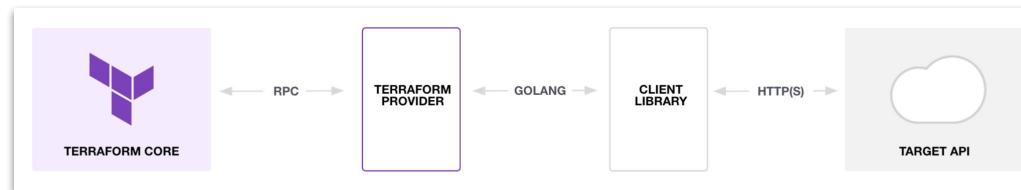
- *Identifiers* can contain letters, digits, underscores (_), and hyphens (-). The first character of an identifier must not be a digit, to avoid ambiguity with literal numbers.
- Comments can be defined by #, //, or /* */.
- Configuration files must always be UTF-8 encoded



Plugin-Based Architecture

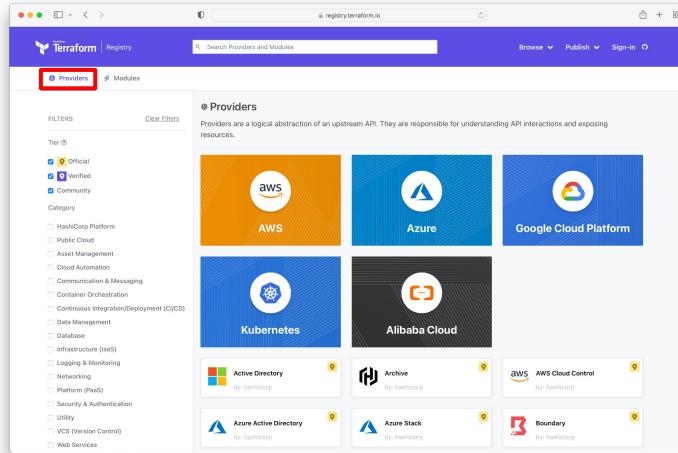
Core binary is small, providers are developed as plugins

- *Terraform Core:* The binary that communicates with plugins to manage infrastructure resources
- *Terraform plugins:* Executable binaries written in Go that communicate with Terraform core over an RPC interface



Exploring the Provider Registry

Central location for provider plugins and their documentation



<https://registry.terraform.io/browse/providers>



Defining Required Provider(s)

Sets up the properties for all providers of the assigned name

The free-form
name of the
provider

```
terraform {  
  required_providers {  
    aws = {  
      source  = "hashicorp/aws"  
      version = "6.25.0"  
    }  
  }  
}
```

The location of the
provider plugin in the
registry in the format
[<HOSTNAME>/] <NA
MESPAC>/<TYPE>

The version (range)
selector for the plugin



Version Constraint & Selection

You can define a concrete version or a version range

Version Value	Meaning
no assigned value attribute	Picks the latest version of the provider available in registry
<code>>= 1.2.1</code>	Great than or equal to the version 1.2.1
<code><= 1.2.1</code>	Less than or equal to the version 1.2.1
<code>~> 1.2.1</code>	Any version in the 1.2.x range up to the next minor version (1.3.0)
<code>>= 1.2.1, <= 1.5.0</code>	Any version between 1.2.1 and 1.5.0



Dependency Lock File

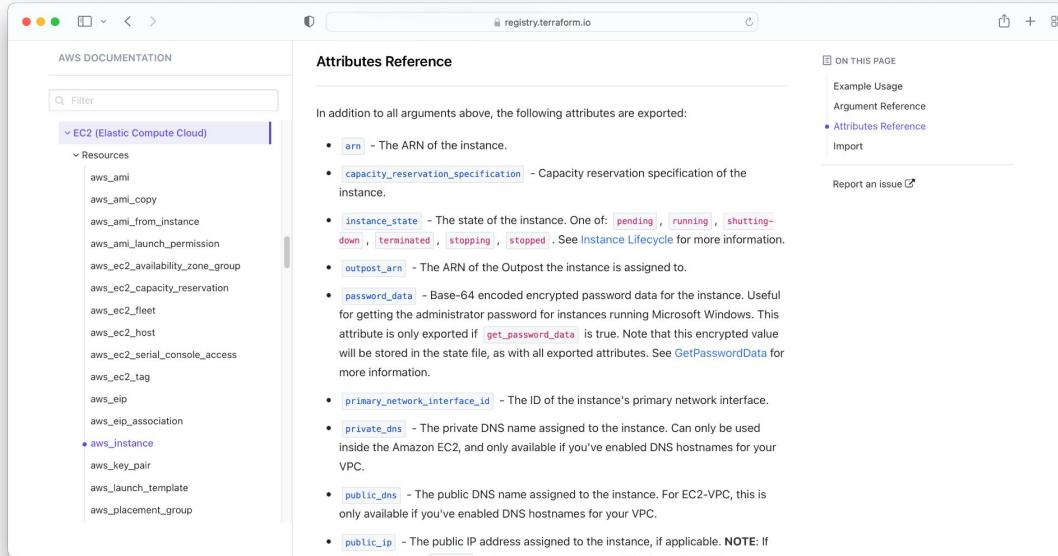
Stores resolved versions + checksum for provider plugins ([docs](#))

- Terraform creates or updates the dependency lock file each time you run the `init` command.
- The lock file is always named `.terraform.lock.hcl` and is meant to be checked into version control.
- Use the `--upgrade` flag with the `init` command to upgrade a provider plugin version.



Provider Documentation

A provider publishes the API (e.g. available attributes) on registry



Configuring Provider(s)

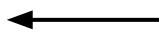
Need to be defined in the root module configuration file

```
provider "aws" {  
  region = "us-east-2"  
}
```



Default provider: Any resource that does not assign a provider explicitly will use this provider

```
provider "aws" {  
  alias = "west"  
  region = "us-west-2"  
}
```



Aliased provider: A resource can refer to this provider explicitly by assigning the provider argument (in this case aws.west)



Local Plugin Cache

Core binary is small, providers are developed as plugins

- Terraform automatically fetches plugins from the registry and stores them in the `.terraform` subdirectory
- Caching is enabled by default and a plugin version already available in the `.terraform` subdirectory will be reused
- Cache directory be configured with the environment variable `TF_PLUGIN_DIR` or the CLI option `--plugin-dir`



Configuring Provider Credentials

Provider exposes attributes for consuming credentials

```
provider "aws" {
  region = "us-east-2"
  access_key = "my-access-key"
  secret_key = "my-secret-key"
}
```



Do you not use this way to provide credentials!
The state file with store credentials in plain text.

```
provider "aws" {
  region = "us-east-2"
  shared_credentials_files = ["path-to-credentials-file"]
  profile = "profile-name"
}
```



Defining Resource(s)

Virtual server in AWS aka EC2 instance

```
resource "aws_instance" "web_server" {
    ami              = "ami-0c55b159cbfafef0"
    instance_type   = "t2.micro"
}
```

Type of EC2 instance which provides a certain amount of CPU, memory, disk space, and network capabilities

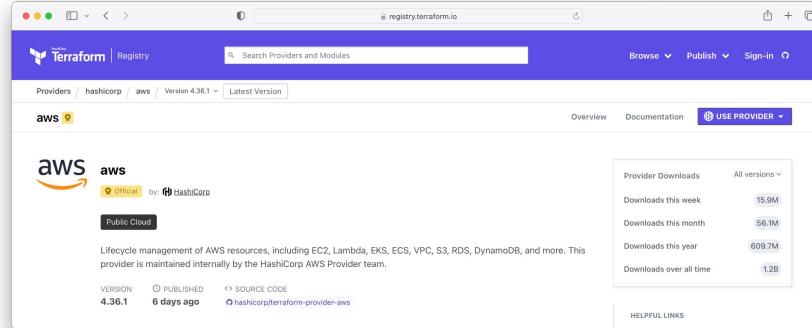
The free or paid Amazon Machine Image (AMI) to run on EC2 instance available via the [EC2 Locator](#)



Provider is Derived from Resource

Defining a provider is optional, Terraform tries to anticipate

```
resource 'aws_instance' "web_server" {}
```



Use the prefix to determine provider on Terraform Registry and download latest version available



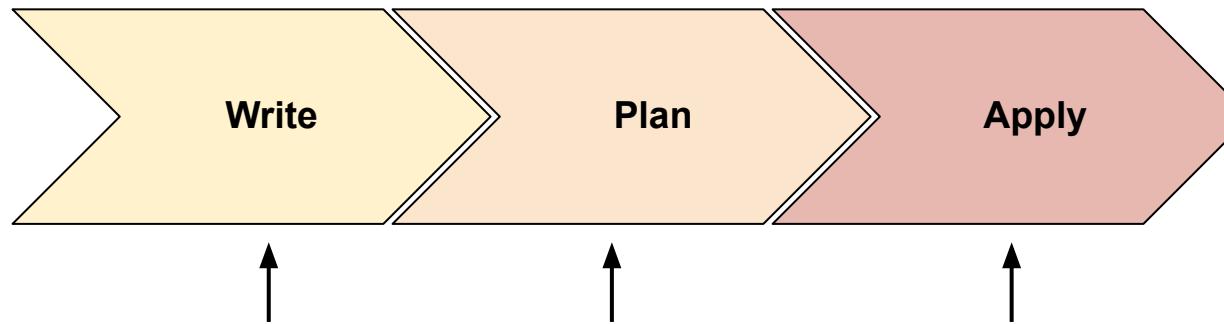
EXERCISE

Implementing a
Simple
Configuration File



Terraform Core Workflow

Three-step approach on a high-level

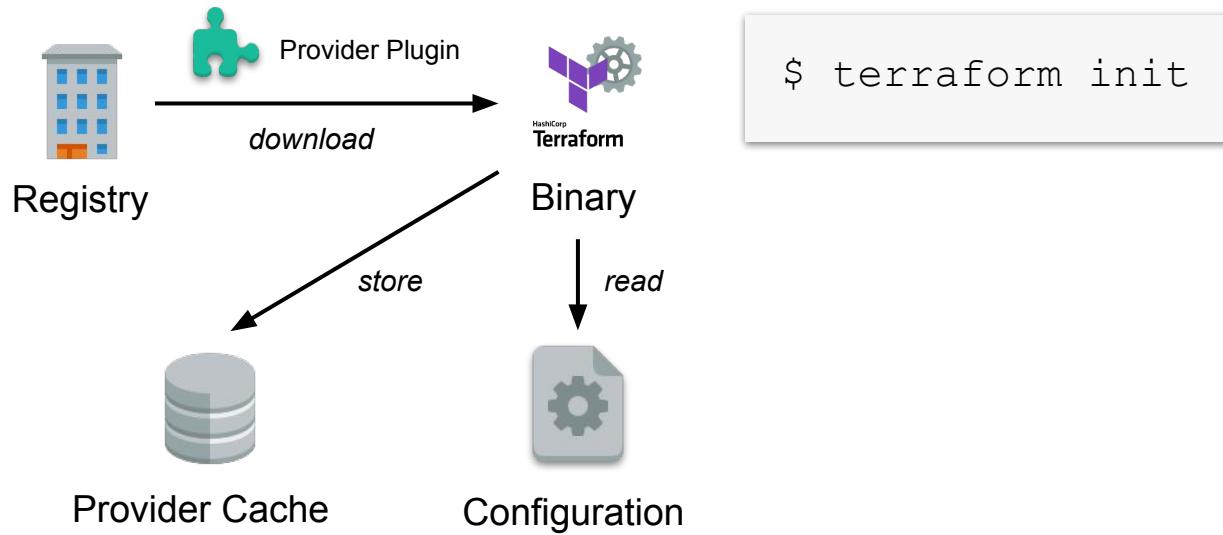


- Write configuration files and run `terraform init` to initialize the configuration to the target environment
- Source code and check it for syntax errors against the target environment
- Terraform syntax can be automatically formatted and validated



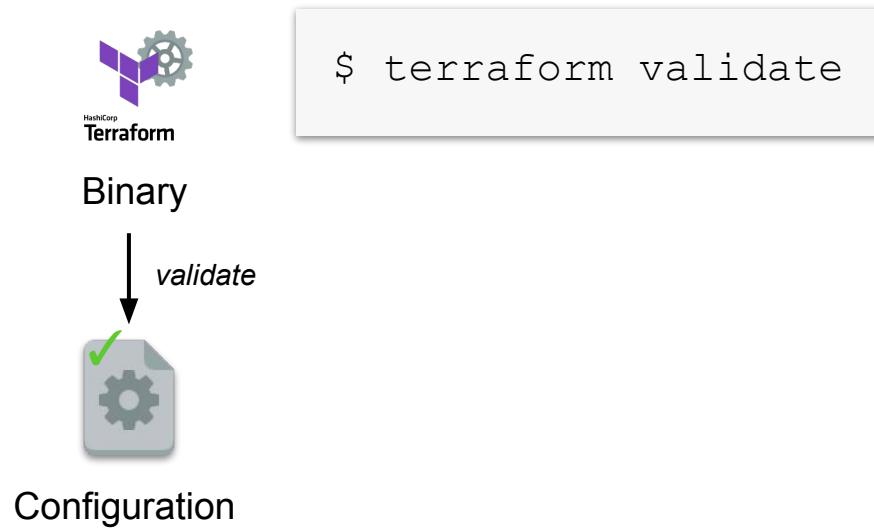
Initialize Working Directory

Prepare working directory, parse config, retrieve provider plugins



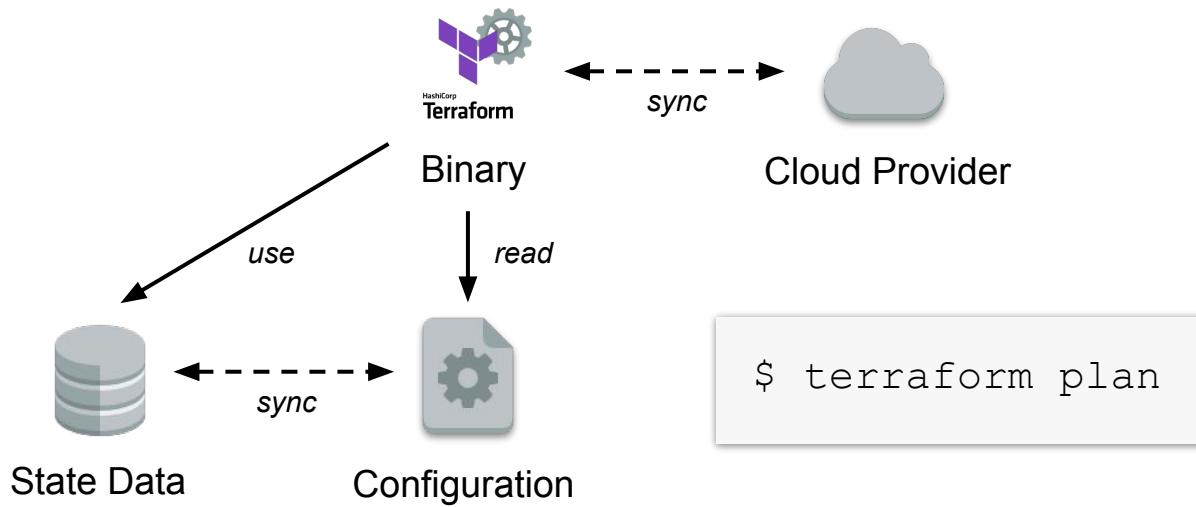
Validate Configuration

Check for syntax errors, doesn't guarantee successful deployment



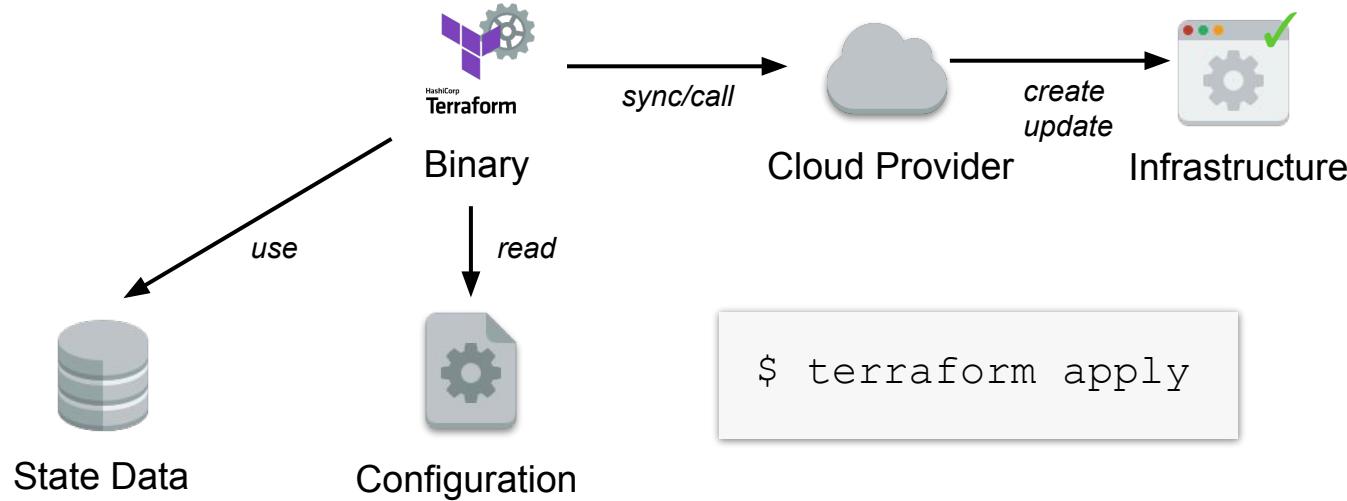
Plan Changes

Determine differences between local state and deployed infra



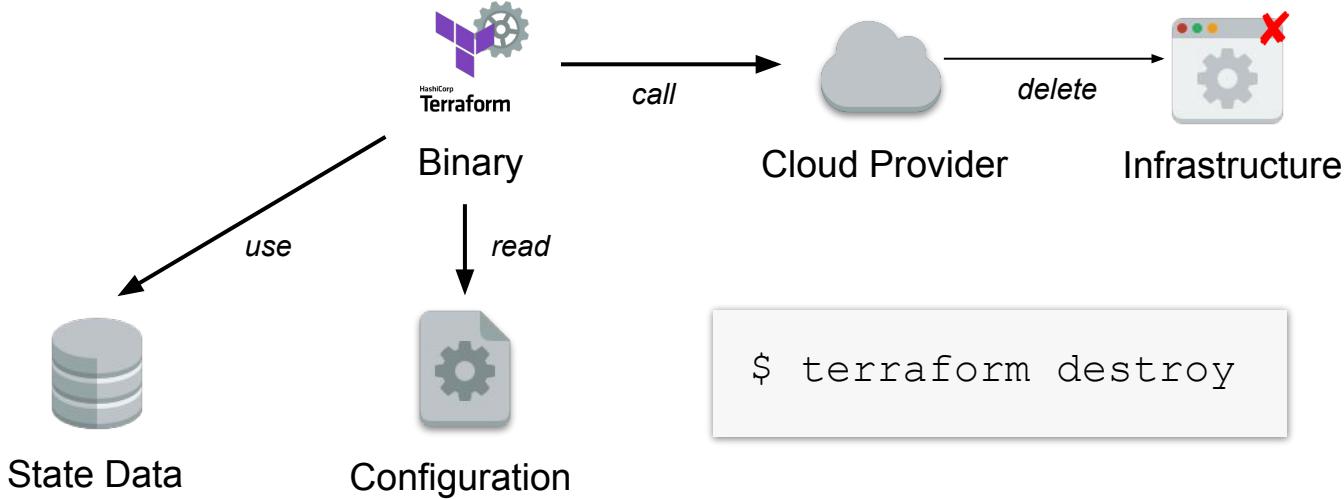
Apply Changes

Deploying the infrastructure or changing its delta



Destroy Infrastructure

Delete infrastructure in target environment based on state data



EXERCISE

Exercising the Core
Terraform Workflow



Q & A



Input Variables, Local/Output Values, and Data Sources

Data Types, Resource Addressing,
Variables/Values

Data Types

A value follows one of the data types below

- Primitive
 - string
 - number
 - bool
- Collection or Complex
 - list/tuple
 - set
 - map/object



Primitive Data Types

Definition very similar to programming/scripting languages

```
# string  
"hello" or "this is an example"
```

```
# number  
25 or 4.632
```

```
# bool  
true or false
```



Collection Data Type Examples

Items in collection can be defined on individual lines for readability

```
# list (or tuple)
["us-west-1a", "us-west-1c"]
```

```
# map (or object)
{name = "Mabel", age = 52}
```



Resource Addressing

Referencing a variable or resource attribute in a different context

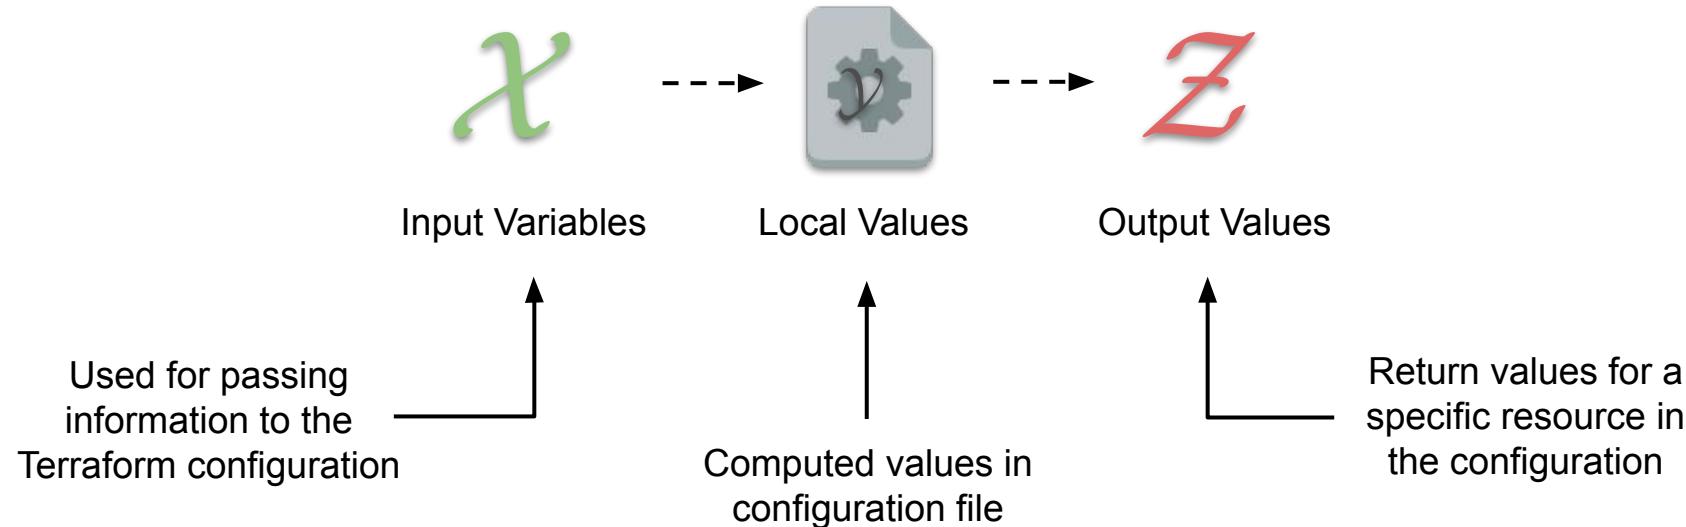
```
resource "aws_instance" "web_server" {
    ami           = "ami-0c55b159cbfafef0"
    instance_type = "t2.micro"
}
```

aws_instance.web_server.ami



Defining Named Values

Concepts for requesting, referencing, and publishing values



Use Cases for Named Values

Understanding when to use which type



Input Variables



“I want to capture values from the CLI when the end user invokes Terraform and use it in my configuration.”



Local Values



“I want to create a variable, assign a value, and reuse it in my configuration similar to a constant.”



Output Values



“I want to render a runtime value in the console output or use it as an input for a resource or module.”



File Naming Conventions

Terraform will automatically resolve those files, but not enforce



Input Variables

variables.tf



Local Values

locals.tf



Output Values

outputs.tf



Defining Input Variables

Configuration file can define 0 to many variables

```
variable "ami_id" {
    type = string
    description = "The AMI identifier to use
                    for EC2 instance."
}

variable "availability_zone_names" {
    type      = list(string)
    default  = ["us-west-2"]
}
```



Input Variable Arguments

All arguments are optional

Argument	Description
default	Value to be used if not provided
type	Accepted type for values assigned to variable
description	End user description that explains purpose and kind of value
validation	Validation rules applied to provided value
sensitive	Obfuscates sensitive information in CLI output
nullable	Defines if assigned value can be <code>null</code>



Referencing Input Variable Values

A variable can be used with the `var.name` notation

```
resource "aws_instance" "example" {
    ami          = var.ami_id           ← string value
    instance_type = var.availability_zone_names[0]
}
```



First element in list



Requesting Variable Input Values

End user has to provide value if no default value has been set

```
$ terraform plan  
var.ami_id ← Variable name  
The AMI identifier to use for EC2 instance.  
  
Enter a value: ami-077ee47512dc6f3ca
```

The diagram illustrates the interaction between Terraform command-line output and user input. A horizontal line with arrows points from the user's input 'Enter a value:' back to the 'var.ami_id' variable definition in the command-line output. Three labels with arrows point to specific parts of the interface: 'Variable name' points to the label 'var.ami_id', 'Description' points to the explanatory text 'The AMI identifier to use for EC2 instance.', and 'End user-provided values' points to the user's input 'Enter a value: ami-077ee47512dc6f3ca'.

Variable name

Description

End user-provided values



Variable Input Values from CLI

-var option when running the plan and apply commands

```
$ terraform apply -var="ami_id=ami-0c55b159cbfafef0"  
...  
$ terraform apply  
-var='availability_zone_names=["us-east-1a", "us-west-1c"]'  
...
```



Variable Input Values from Files

-var-file option can point to a file containing input variables

```
$ terraform apply -var-file="runtime-var.tfvars"
```

```
ami_id = "ami-0c55b159cbfafef0"
availability_zone_names = [
  "us-east-1a",
  "us-west-1c"
]
```

runtime-var.tfvars



Auto-Loading of Variable Files

Standard naming convention will apply

- Files named exactly `terraform.tfvars` or `terraform.tfvars.json`
- Any files with names ending in `.auto.tfvars` or `.auto.tfvars.json`



Validating Input Variables

Built-in functions can be used to implement validation logic

```
variable "ami_id" {  
    ...  
    validation {  
        condition = length(var.ami_id) > 4 &&  
                    substr(var.ami_id, 0, 4) == "ami-"  
        error_message = "The image_id value must be a valid  
                        AMI id, starting with \"ami-\"."  
    }  
}
```



Multiple Validation Rules

Reports all failures of the whole stack of validation blocks

```
variable "environment" {  
  ...  
  validation {  
    condition = contains(["dev", "staging", "prod"], var.environment)  
    error_message = "Environment must be dev, staging, or prod."  
  }  
  validation {  
    condition = length(var.environment) >= 3  
    error_message = "Environment name must be at least 3 characters."  
  }  
}
```



EXERCISE

Input Variable
Definition and
Consumption



Secrets in State

Terraform stores all values in state, including secrets

```
resource "aws_db_instance" "main" {  
    password = var.db_password ← Stored in state  
}
```

- This is a major concern in enterprise Terraform
- State files contain passwords in plain text
- Even with S3 encryption at rest, anyone with state access can read secrets



Injecting Secrets as Inputs

Input variables will be stored in the state in plain text

```
$ export TF_VAR_third_party_pwd=s3cr3t  
$ export TF_VAR_api_key=a8an23sdf023xmkdd
```

main.tf

```
# use variables  
var.third_party_pwd  
var.api_key
```



Ephemeral Values

Never persisted to state or plan files

- **Ephemeral input variables:** Sensitive inputs not stored
- **Ephemeral resources:** Temporary data that exists only during operations
- **Write-only arguments:** Resource attributes not persisted (Terraform 1.11+)



Ephemeral Resources

Generate or fetch temporary data

```
ephemeral "random_password" "db" {  
    length          = 16  
    special         = true  
    override_special = "!#$%&* ()-_=_+=[]{}<>:_?"  
}
```

Executed during every plan and apply. Its values are not stored in state.



Ephemeral Resources Lifecycle

No data persistence of value

1. **Opening:** When required, the resource is "opened" (e.g., a secret fetched from a vault), generating a temporary value.
2. **Renewing (Optional):** If the resource's lease or validity is short, the provider can periodically "renew" it to extend its lifespan during the operation.
3. **Closing:** Once all dependent actions finish, the resource's lease is revoked, ensuring cleanup and preventing data persistence.



Write-Only Arguments

Accept values but never store them (Terraform 1.10+)

```
resource "aws_db_instance" "main" {
    engine          = "postgres"
    instance_class = "db.t3.micro"
    username        = "admin"

    password_wo      = ephemeral.random_password.db.result
    password_wo_version = 1
}
```



The `_wo_version` argument tracks when to update the value.



Write-Only Version Management

Triggers the update

- The version is just an integer - increment to trigger updates
- Some teams use a timestamp or hash instead
- Pattern for rotation: increment version in a PR, merge to trigger rotation
- Consider storing version in a separate variable for easier management



Supported Write-Only Resources

As of Terraform 1.12, write-only arguments are available in:

Provider	Resources
AWS	<code>aws_db_instance</code> , <code>aws_secretsmanager_secret_version</code>
Azure	<code>azurerm_mssql_server</code> , <code>azurerm_key_vault_secret</code>
Google	<code>google_sql_user</code> , <code>google_secret_manager_secret_version</code>

Check provider documentation for the latest supported resources.



EXERCISE

Ephemeral Value
Definition and
Consumption



Defining Local Values

Supports reusability of the same value in a configuration

```
locals {
    some_other = var.from_cli
    default_tags = {
        Organization = "O'Reilly"
        Owner = "Benjamin Muschko"
    }
}
```

Parsed value from
input variable



Referencing Local Values

A local value can be used with the `local.name` notation

```
provider "aws" {
  region = "us-east-2"
  default_tags {
    tags = local.default_tags
  }
}
```



Assigns the local value from
the locals definition block



EXERCISE

Local Value
Definition and
Consumption



Defining Output Values

Make information about infrastructure available on CLI

```
output "example_ip_address" {
    value = aws_instance.example.private_ip
    description = "The private IP address of the
                   main server instance."
}
```

References the
private IP address
attribute of the
AWS instance
named example



Output Values in CLI

Outputs can be queried from the state database (if populated)

```
$ terraform output
Warning: No outputs found
The state file either has no outputs defined, or all the
defined outputs are empty.

$ terraform apply
example_ip_address = "..."

$ terraform output
example_ip_address = "..."
```



EXERCISE

Output Value
Definition and
Rendering



Defining Data Sources

Get information about resources external to Terraform

```
# define data source
data "github_repository_pull_requests" "pull_requests" {
    base_repository = "example-repository"
    base_ref = "main"
    state = "open"
}

# Reference data
data.github_repository_pull_requests.pull_requests.results
```

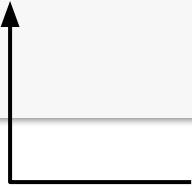


Check Block

Validate infrastructure without blocking operations

```
check "website_health" {
  data "http" "app" {
    url = "https://${aws_lb.main.dns_name}/health"
  }

  assert {
    condition      = data.http.app.status_code == 200
    error_message = "Application health check failed."
  }
}
```



Failed checks produce warnings, not errors.



EXERCISE

Defining and
Consuming a Data
Source



Resource Dependencies

Implicit/explicit dependencies, Lifecycle block

Implicit Dependencies

Detects dependencies when one resource references another

```
resource "aws_vpc" "main" {
  cidr_block = "10.0.0.0/16"
}

resource "aws_subnet" "web" {
  vpc_id      = aws_vpc.main.id
  cidr_block = "10.0.1.0/24"
}
```

Creates an implicit
dependency



Implicit Dependencies

Preferred approach - let Terraform figure it out

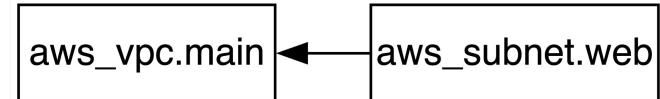
- Terraform builds a dependency graph and creates resources in the correct order
- The dependency graph is visible with `terraform graph` command



Rendered Graph

Provides a GraphViz representation

```
$ terraform graph
digraph G {
    rankdir = "RL";
    node [shape = rect, fontname = "sans-serif"];
    "aws_subnet.web" [label="aws_subnet.web"];
    "aws_vpc.main" [label="aws_vpc.main"];
    "aws_subnet.web" -> "aws_vpc.main";
}
```



<https://dreampuf.github.io/GraphvizOnline/>



When the approach doesn't work

Sometimes Terraform can't detect dependencies

- IAM policies that must exist before a resource can use them
- Security groups referenced by name instead of ID
- Resources that depend on side effects of other resources
- External systems that need time to propagate changes



Explicit Dependencies

Instance could launch before the IAM policy is attached

```
resource "aws_iam_role_policy" "s3_access" {
  name      = "s3-access"
  role      = aws_iam_role.app.id
  policy    = jsonencode({...})
}

resource "aws_instance" "app" {
  ami           = "ami-0c55b159cbfafef0"
  instance_type = "t3.micro"
  depends_on    = [aws_iam_role_policy.s3_access]
}
```

Terraform
cannot infer
dependencies



When to use `depends_on`

Valid reasons for declaring explicit dependencies

- Hidden dependencies on IAM policies/roles
- Dependencies on resources accessed by name
- Cross-account or cross-region dependencies
- When order of creation matters for external reasons



When not to use `depends_on`

In most cases you won't need explicit dependency declarations

- When a direct reference creates an implicit dependency
- As a workaround for timing issues (use `time_sleep` instead)
- For data sources (can cause circular dependencies)



Lifecycle Block

Overrides Terraform's default resource management behavior

- The `lifecycle` block is one of Terraform's most powerful features
- It's a nested block that goes inside the `resource` block
- Available for all resource types - it's a meta-argument
- Customizes how Terraform manages the resource during create, update, and destroy operations



Core Lifecycle Arguments

Modify the default behavior of resource management

Argument	Runtime Behavior
<code>create_before_destroy</code>	Creates a new replacement resource before deleting the old one to ensure zero downtime.
<code>prevent_destroy</code>	Prevents Terraform from deleting a resource; any plan to destroy it will result in an error.
<code>ignore_changes</code>	Stops Terraform from reverting manual or external updates to specific resource attributes.
<code>replace_triggered_by</code>	Forces resource replacement when a specified dependency or attribute changes elsewhere.



create_before_destroy

Implementing a zero-downtime control

- **Default behavior:** Terraform destroys the old resource, then creates the new one.
- **Use case:** Zero-downtime deployments where the resource must remain available.
- With `create_before_destroy = true`: Terraform creates the new resource first, then destroys the old one.



create_before_destroy

New EC2 instance will be created before old one is destroyed

```
resource "aws_instance" "web" {
    ami           = var.ami_id
    instance_type = "t3.micro"

    lifecycle {
        create_before_destroy = true
    }
}
```



Usage Considerations

May require additional customization

- Resources with unique name requirements may conflict
- May require random suffixes or name prefixes
- Not all resources support having two instances simultaneously



Preconditions

Validate things outside the resource itself

```
resource "aws_instance" "web" {
  ...
  lifecycle {
    precondition {
      condition = data.aws_ami.selected.architecture == "x86_64"
      error_message = "The AMI must use x86_64 architecture."
    }
  }
}
```



Postconditions

Validate guarantees after Terraform creates a resource

```
resource "aws_instance" "web" {
  ...
  lifecycle {
    postcondition {
      condition = self.public_ip != ""
      error_message = "Instance must have a public IP address."
    }
  }
}
```



Preconditions vs Postconditions

Both block the operation on failure

Aspect	Precondition	Postcondition
When evaluated	Before resource creation	After resource creation
Can use <code>self</code>	No	Yes
Purpose	Validate assumptions	Validate guarantees
Typical use	Check inputs, data sources	Verify outputs, attributes



Q & A



Iteration Control, Functions, and Expressions

Loops, Dynamic Blocks, Built-In Functions

Looping Constructs

Terraform provides looping syntax for different use cases

- count argument: Assigns an integer that determines # of loops
- for_each argument: Iterates over each element of a map or set
- for expression: Used to transform elements in a map or set
- dynamic blocks: Similar to a traditional for-loop but uses the dynamic nested block syntax



Count Argument

Count is a simple integer value assignment

```
resource "aws_instance" "app_server" {
  count = 3
  ami   = "ami-0c55b159cbfafef0"
  instance_type = "t2.micro"
  tags = {
    Name = "app_server_${count.index}"
  }
}
```



Current index of loop
can used as variable



For Each Argument

Iteration element can be referenced by key and value attribute

```
resource "azurerm_resource_group" "rg" {
  for_each = {
    a_group = "eastus"
    another_group = "westus2"
  }
  name      = each.key
  location  = each.value
}
```

map data type with 2 entries

For a set data type the key
is the same as value



For Expression

Applies an expression to each element to transform it

```
output "all_tags" {
    value = {for key, value in var.tags :
              upper(key) => upper(value) }
}
```

←
Iterate over
elements in a
map

↑
Transforms key and
value with function
and outputs a map



Dynamic Block

Helps with implementing DRY for more elaborate configurations

```
resource "aws_security_group" "security" {
  ...
  dynamic "ingress" {
    for_each = local.ingress_rules
    ingress {
      description = ingress.value.description
      from_port   = ingress.value.port
      to_port     = ingress.value.port
      protocol    = "tcp"
      cidr_blocks = ["0.0.0.0/0"]
    }
  }
}
```

Local variable pointing to a set of tuples with port numbers and descriptions

Configure the whole “block” for each iteration the set



EXERCISE

Declaring a Loop
Sourced from an
Input Variable



Built-In Functions

Functionality for transforming, combining, and formatting data

- Grouped into higher-level categories e.g. collections, type conversion, date and time ([docs](#))
- Syntax in code: `function_name(argument, ...)`
- Functions can be combined with each other by nesting them
- Terraform provides a console to try out functions against state data



Function Examples

Available functions are vast, suitable for most common use cases

```
# return largest value in a set
max(22, 5, 99)
> 99

# base64-encode a string value
base64encode( "Hello World")
> SGVsbG8gV29ybGQK

# extract a substring from a given string
substr("hello world", 1, 4)
> ello
```



Terraform Console

Test functions against the current state data (and locks it)

```
$ terraform console
> var.instance_count
3
> range(0, var.instance_count)
tolist([
  0,
  1,
  2,
])
```

Get the value of the input variable `instance_count`

Execute the `range` function against the variable



EXERCISE

Using Built-In
Functions to
Transform Data



Interacting with Modules

Module Structure and Syntax, Inputs & Outputs

What is a Module?

Reusable functionality in the world of Terraform

- Usually defines inputs, resources/data sources, and outputs (though they are all optional)
- Modules are versioned (standardized on semantic versioning)
- The `init` command will download remote modules and store them in the `.terraform` directory



Module Sources

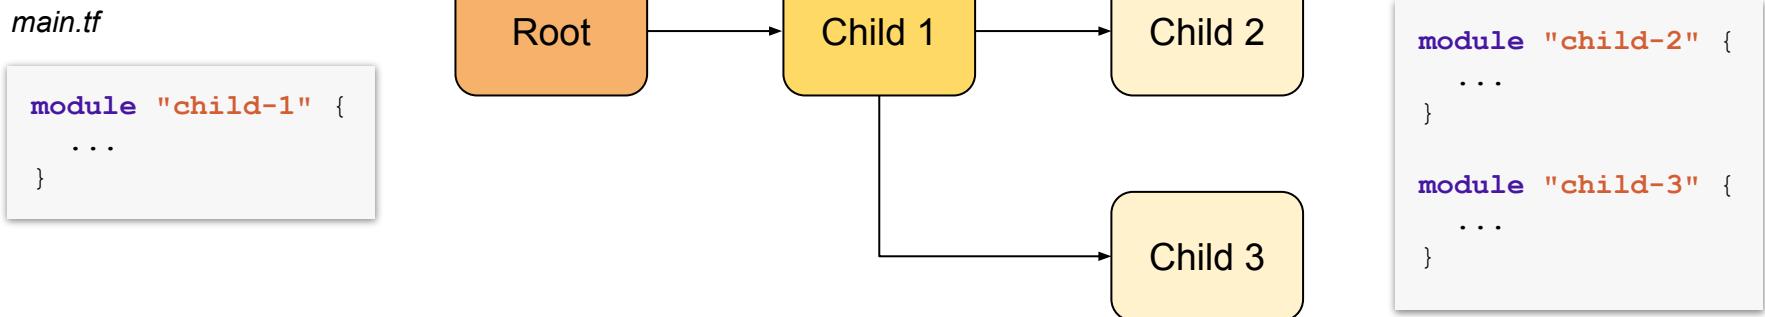
A module can be sourced from different locations

- A local path to a directory present on disk
- A registry like the [Terraform registry](#)
- A Git repository e.g. GitHub or BitBucket
- A HTTP(S) URL that follows a provider protocol



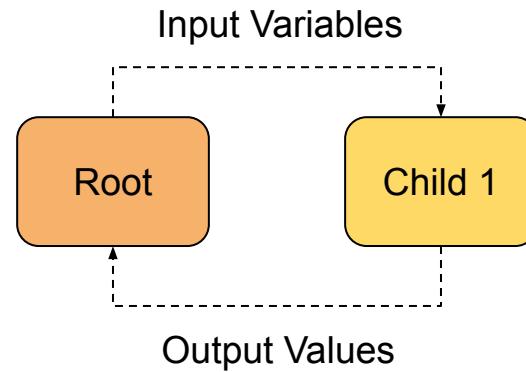
Root and Child Modules

Modules can be consumed on any level of the hierarchy



Module Data Exchange

Modules only communicate via inputs and outputs



Module Example

Organized in a subdirectory, and defined by .tf files

modules/eks/main.tf

```
variable "cluster_name" {
    type = string
}

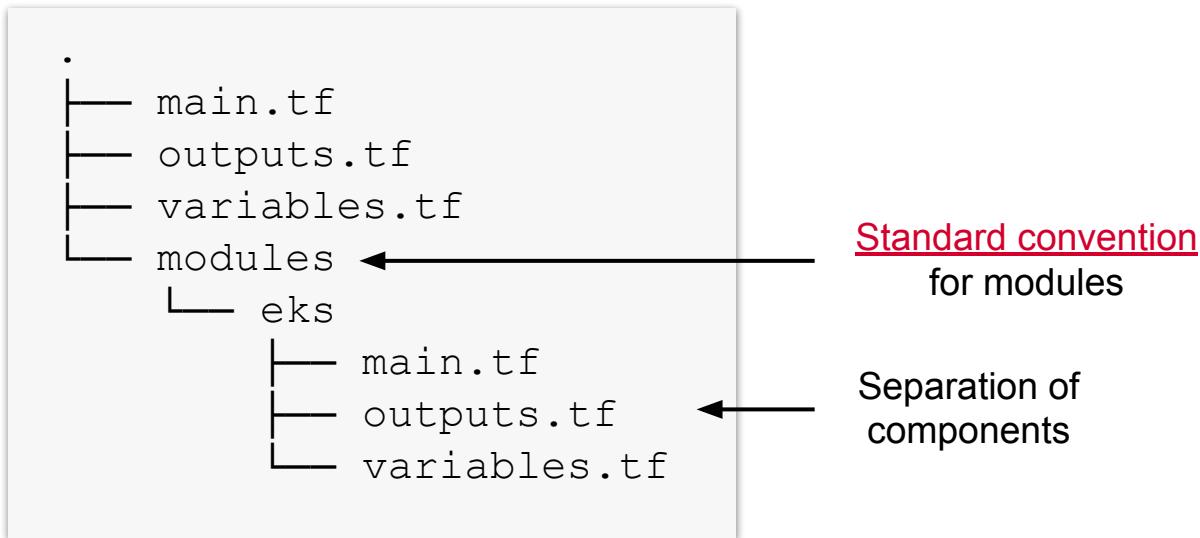
resource "aws_eks_cluster" "eks_example" {
    name = "${var.cluster_name}-eks"
    ...
}

output "endpoint" {
    value = aws_eks_cluster.eks_example.endpoint
}
```



Module Structure

Organized in a subdirectory, and defined by .tf files



Consuming a Local Module

References by source, no version needs to be defined

main.tf

```
module "aws_eks_cluster" {  
  source = "./modules/eks"  
  cluster_name = "my_cluster"  
}  
  
resource "..." {  
  endpoint =  
    module.aws_eks_cluster.endpoint  
}
```

Location of module in directory structure

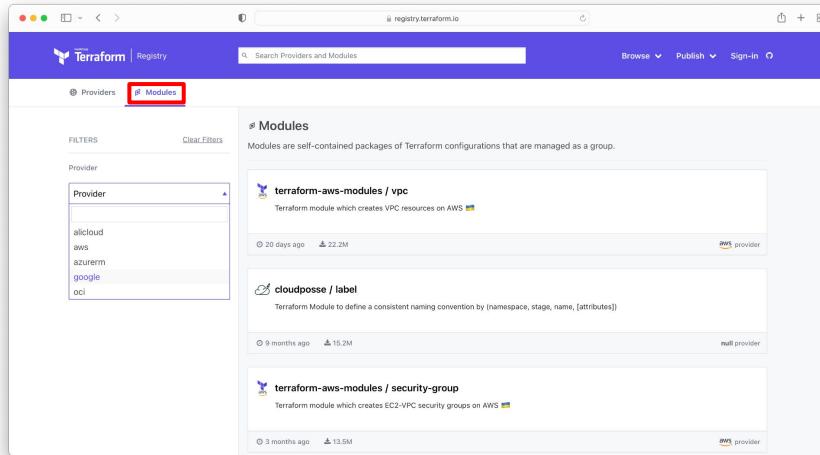
Providing the value for the input variable

Accessing the output value from local module



Exploring the Module Registry

Similar to providers, modules are available Terraform registry



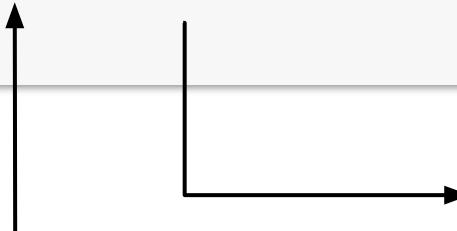
<https://registry.terraform.io/browse/modules>



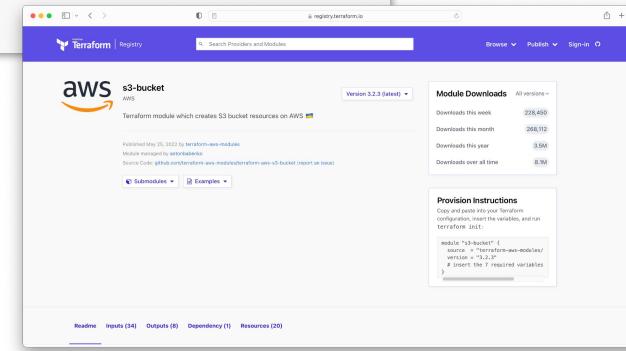
Consuming a Public Module

Define source and version

```
module "s3-bucket" {  
  source = "terraform-aws-modules/s3-bucket/aws"  
  version = "3.2.3"  
}
```



Follows the same
conventions and notation as
provider versions



EXERCISE

Implementing and
Using a Custom
Module



Q & A

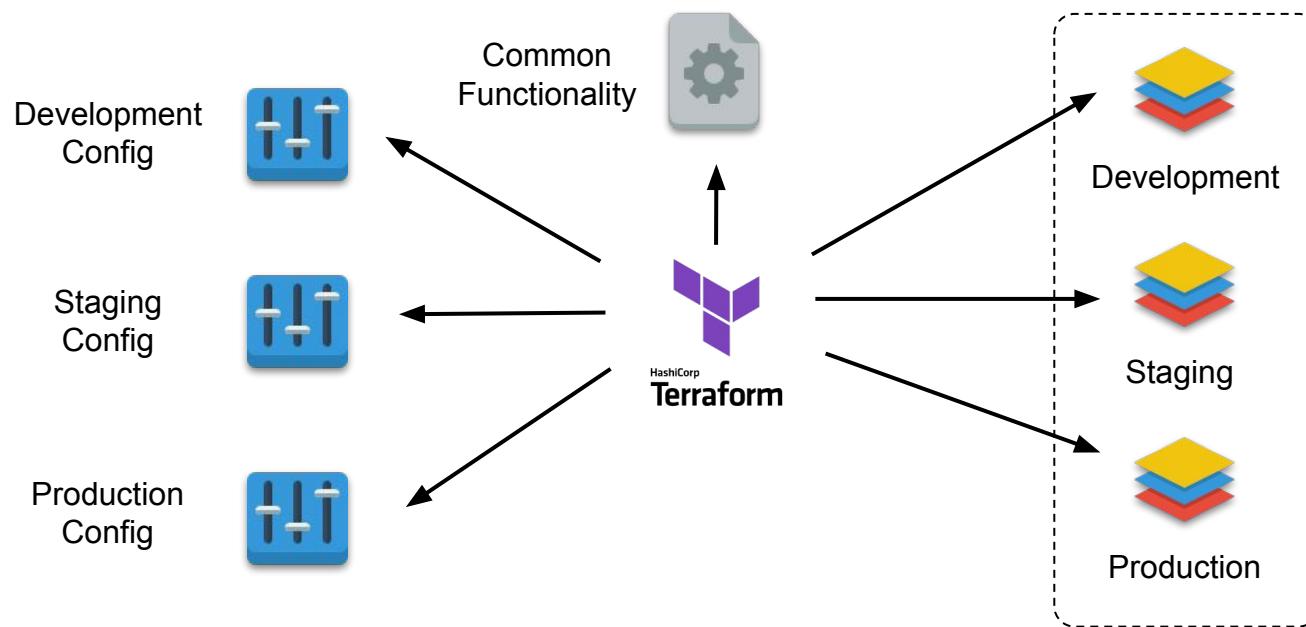


Managing Multiple Runtime Environments

Workspaces, Commands, Variable Values,
Separated State

Multiple Environments

Commonalities and differences



Handling Environments

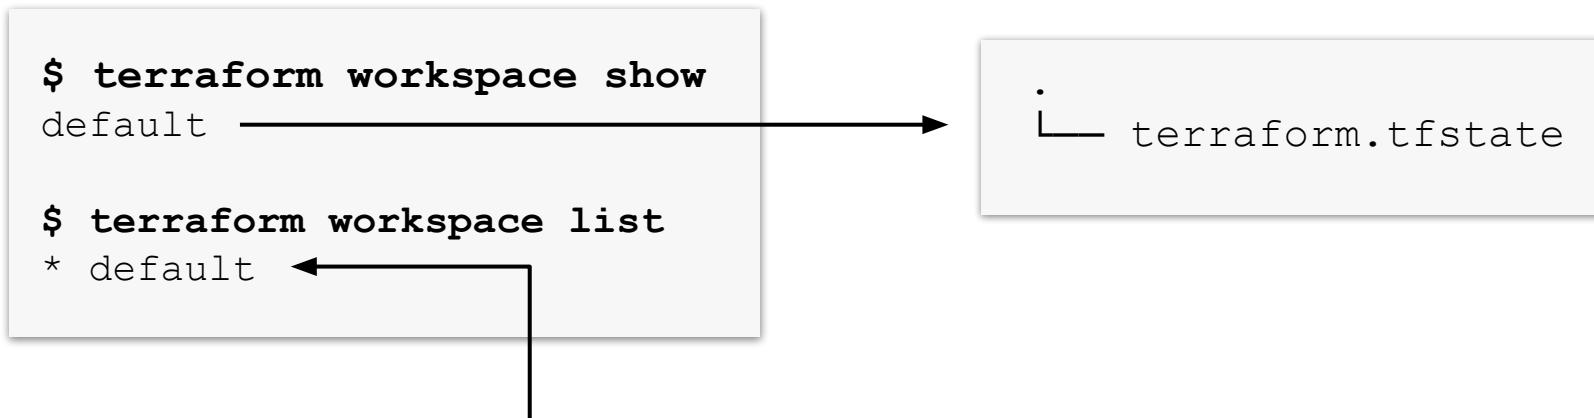
A workspace manages requirements

- Shared configuration, different input variable values per context
- Unique credentials per environment to decrease security risks
- Holding unique state per environment
- Increased maintenance effort and potential complexity to end users
- HashiCorp discourages use of workspaces and offers alternatives



The Default Workspace

Cannot be deleted and uses state file in root directory



Marks the selected
workspace with an
asterisk character



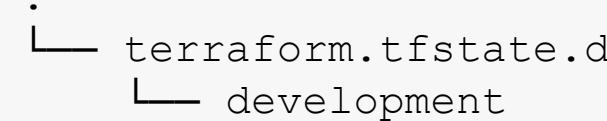
Creating a New Workspace

Built-in command that creates a dedicated state folder

```
$ terraform workspace new development
```

```
Created and switched to workspace "development"!
```

Every workspace
has its own folder



Workspace Variable Values

Per-workspace configuration needs to be resolvable



Development



Staging



Production

Variable	Value
instance_type	t2.nano
instance_count	1

Variable	Value
instance_type	t2.micro
instance_count	2

Variable	Value
instance_type	t2.small
instance_count	4



Defining a Workspace Value

Specify a variable of type map in locals.tf file

```
locals {  
    ec2_instance_type = {  
        development = "t2.nano"  
        staging = "t2.micro"  
        production = "t2.small"  
    }  
  
    ec2_instance_count = {  
        ...  
    }  
}
```

← Assign key-value pairs per environment

← Repeat for every workspace-specific variable

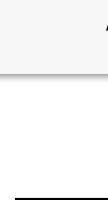


Resolving a Workspace Value

Use the variable `terraform.workspace` to select value

```
resource "aws_instance" "web_server" {  
    ami = "ami-0c55b159cbfafef0"  
    instance_type = local.ec2_instance_type[terraform.workspace]  
    instance_count = local.ec2_instance_count[terraform.workspace]  
}
```

Selects value in
corresponding map
variable



Switching Between Workspaces

Use `select` command to target a different environment

```
$ terraform workspace list
  default
  development
* production
  staging

$ terraform workspace select staging
Switched to workspace "staging".

$ terraform workspace show
staging
```



Deleting a Workspace

Select different workspace first, then run delete command

```
$ terraform workspace select development  
Switched to workspace "development".
```

```
$ terraform workspace delete staging  
Deleted workspace "staging"!
```

Removes workspace
folder + state file

```
.  
└── terraform.tfstate.d  
    ├── development  
    └── production
```



EXERCISE

Managing Multiple
Environments With
Workspaces



Implementing and Maintaining State

Local vs. Remote Backend, State Locking, Secret Management

Holding State

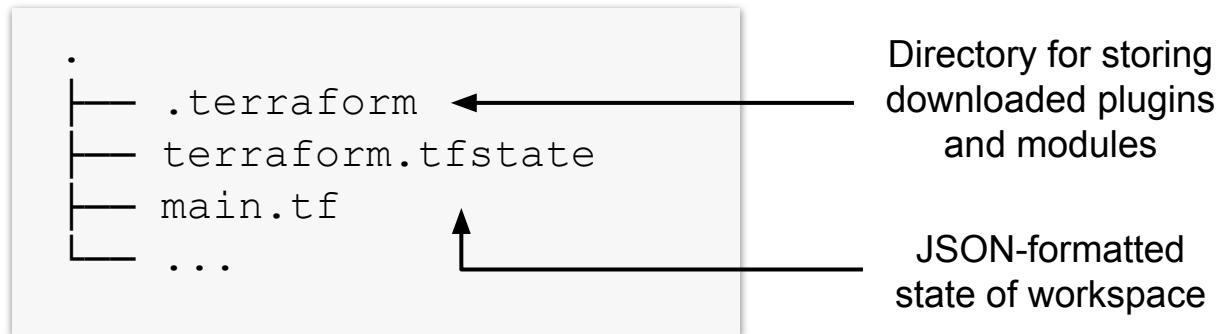
Allows for comparing desired with current infrastructure definition

- Terraform only creates/updates what has actually changed
- Enabler for improving performance
- The commands `plan` and `apply` interact with the state
- Local or remote backend for different use cases



The Default Local Backend

Stored in the `terraform.tfstate` file in workspace directory



Configuring the Local Backend

Default file & location can be changed by CLI or configuration

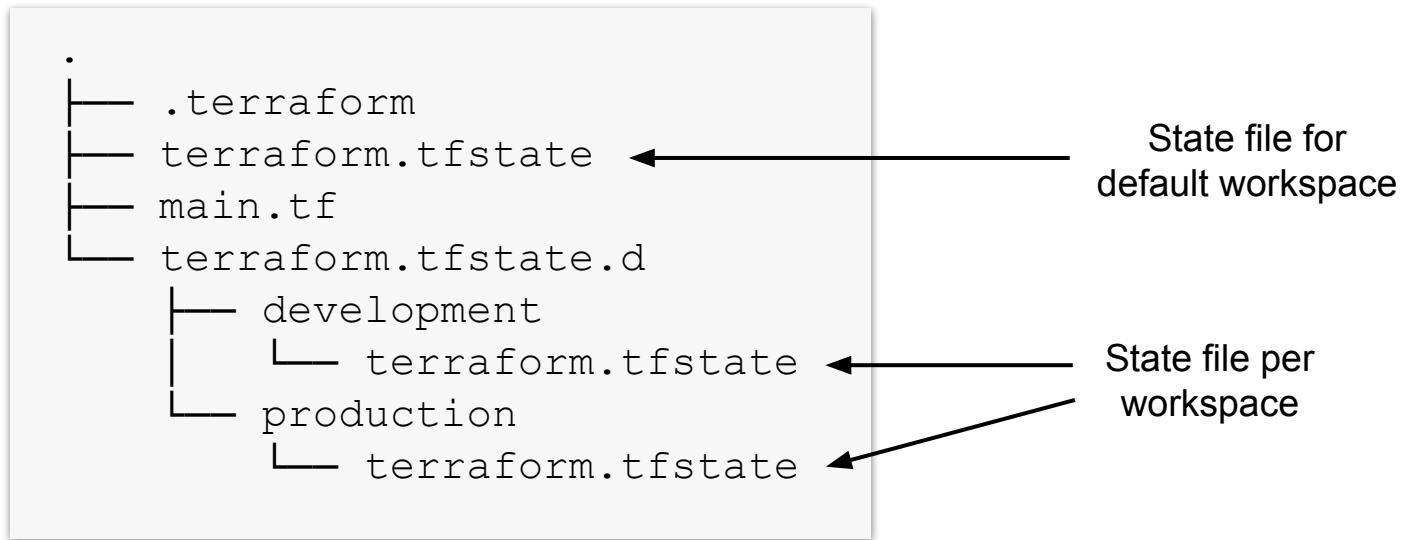
```
terraform {
  backend "local" {
    path = "relative/path/to/terraform.tfstate"
  }
}
```

For ad-hoc tasks, you can use the `-state` command line option



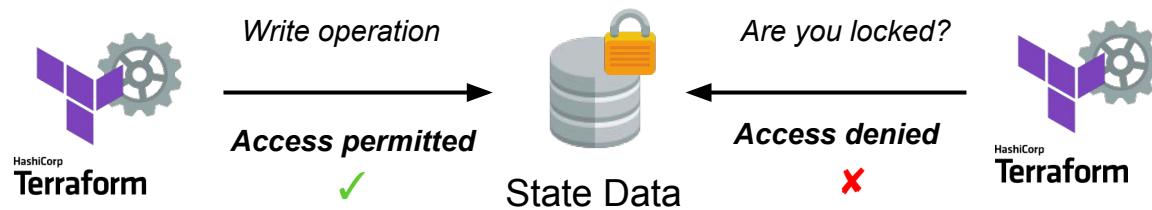
Local State File Per Workspace

State per environment is separated by subdirectory



State Locking

Ensures that state does not corrupted by concurrent access



Lock management can be circumvented using the `-lock=false`.

Caution!



Remote Backends

Optimized for access by multiple users

- Never stores state data on local disk, only loaded into memory and flushed after the operation
- Locking is established on remote backend implementation
- Terraform currently supports different backend [implementations](#)
- Enhanced backends (e.g. HCP Terraform and Enterprise) run Terraform operation on remote service



Remote Backend Authentication

Credentials shouldn't be configured in .tf file

```
terraform {
  backend "remote" {
    hostname = "app.terraform.io"
    organization = "company"

    workspaces {
      name = "my-app-prod"
    }
  }
}
```

Credentials should be provided as environment variables to avoid sending them as plain text over the wire or committing them to version control



Partial Backend Configuration

Information can be provided from the CLI with `init` command

```
terraform {  
  backend "remote" {  
    hostname = "app.terraform.io"  
    organization = "company"  
  
    workspaces {  
      name = "my-app-prod"  
    }  
  }  
}
```

```
$ terraform init  
-backend-config=organization=company
```



HCP Terraform Configuration

Requires a specific notation ([docs](#))

```
terraform {
  cloud {
    organization = "automated-ascent"
    hostname = "app.terraform.io"

    workspaces {
      tags = ["my-app-prod"]
    }
  }
}
```



Secret Management & State

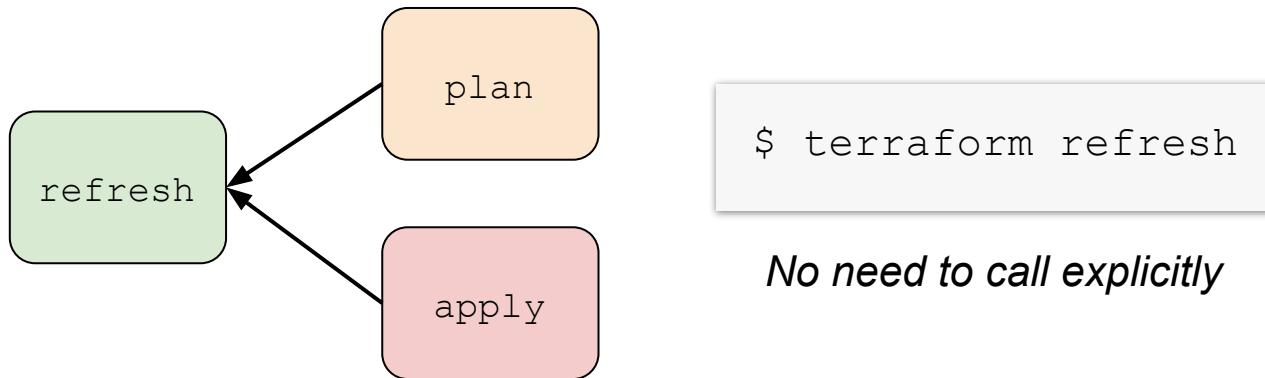
State data is not encrypted and shouldn't store sensitive data

- Sensitive data can be provided with environment variables for some remote backend implementations
- Enhanced backends send and store state information in encrypted form
- Credential-management platforms like [HashiCorp Vault](#) can be integrated as a viable alternative for storing and providing credentials



Managing Resource Drift

Ensures that state synchronizes with deployed infrastructure



The `plan` and `apply` command automatically refresh the state.



Example Refresh Output

Calculated differences between state and cloud are rendered

```
$ terraform plan
Refreshing Terraform state in-memory prior to plan...
The refreshed state will be used to calculate this plan, but will not be
persisted to local or remote state storage.

aws_instance.example: Refreshing state... (ID: i-011a9893eff09ede1)
```

```
-----
An execution plan has been generated and is shown below.
Resource actions are indicated with the following symbols:
~ update in-place
```

```
Terraform will perform the following actions:
  ~ aws_instance.example
    tags.drift_example: "v2" => "v1"
```

Output will list actions
when `apply` command
is performed



Resource Drift Output Symbols

Refresh output shoulds differences and their actions up apply

Symbol	Meaning
+	Create
-	Destroy
-/+	Replace (destroy and then create, or vice-versa if create-before-destroy is used)
~	Update in-place
<=	Read (only applies to data sources)



Deprecation of refresh Command

The command will be removed in a future Terraform version

- The refresh command means the same apply -refresh-only -auto-approve. Therefore, the state will be overwritten. You do have a backup file with the file extension .backup.
- Invalid provider credentials or wrong credentials for a region can lead to an incorrect state.
- It's recommended to run plan/apply with -refresh-only option. The apply explicitly asks for confirmation before overwriting the state file. More information on this [blog post](#).



EXERCISE

Managing Resource
Drift



Advanced Terraform Workflows

fmt, validate, import, state, Enabling verbose
logging

Validating Configuration Files

Check for syntactical errors in configuration files of a directory

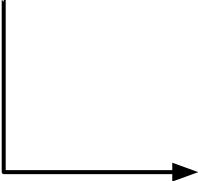
- Functionality provided by the validate command
- Does not consider provided variables or state
- Does not make a guarantee on successful deployment
- Should be run before plan/apply commands and is a good fit for CI/CD pipelines



Validation Example

Renders errors on terminal output

```
instance_type = ["t2.micro"]
```



```
$ terraform validate
Error: Incorrect attribute value type
on main.tf line 16, in resource "aws_instance" "app_server":
16:   instance_type = ["t2.nano"]

Inappropriate value for attribute "instance_type": string required.
```



Formatting Configuration Files

Align configuration with standard style conventions

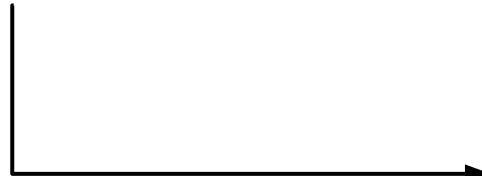
- Functionality provided by the `fmt` command
- Applies to `.tf` and `.tfvars` files in the current directory, use `-recursive` to apply formatting to all subdirectories
- Other CLI options are available: `-list=false`, `-write=false`, `-diff`, `-check`
- Should be performed before committing to version control



Reformatting Example

Aligns multiple attribute assignments for a resource

```
ami = "ami-0c55b159cbfafef0"  
instance_type = "t2.micro"
```



```
ami          = "ami-0c55b159cbfafef0"  
instance_type = "t2.micro"
```



EXERCISE

Validating and
Formatting
Configuration Files



Importing Existing Infrastructure

Bring manually created infrastructure under Terraform's control

- Functionality provided by the `import` command
- Loads supported resources into your Terraform workspace's state
- Does not automatically generate the configuration to manage the infrastructure
- Importing infrastructure is a [multi-step process](#)



Identifying the Resource

Every resource to be imported needs to have an identifier

Instances (1) Info											
		Actions Launch instances									
<input type="button" value="Search"/> <input type="text" value=""/>		Connect Instance state Actions Launch instances									
<input type="button" value="Instance state = running"/> X Clear filters											
Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public IPv4 IP	Elastic IP		
-	i-094fd5c99f77c7afe	Running	t2.nano	2/2 checks passed	No alarms	+ us-west-2c	ec2-54-214-82-80.us-w...	54.214.82.80	-		



For an EC2 instance,
the resource identifier
is the *instance ID*



Adding a Resource Config Block

The `import` command will attach state to configuration

```
resource "aws_instance" "web_server" {}
```

Do not add any specific
configuration yet



Running the Import Command

Provide the configuration ID + resource ID

```
$ terraform import aws_instance.web_server i-094fd5c99f77c7afe
```



Configuration ID



Resource ID



Check Imported State

Render the resource state that could be imported

```
$ terraform show
```



Manually Add Configuration

Fill in missing information that couldn't be derived

```
resource "aws_instance" "web_server" {  
    ...  
}
```

Add the
configuration of
the resource



Plan and Apply Changes

Iteratively plan, add configuration, and finally apply

```
$ terraform plan  
$ terraform apply
```



Import Block

Terraform 1.5+ makes things easier ([docs](#))

```
import {
  to = aws_instance.example
  id = "i-abcd1234"
}

resource "aws_instance" "example" {
  name = "backend"
}
```

```
$ terraform plan
```



Interacting with Terraform State

Fine-grained interactions with state backend

- Manual updates to the state document, local or remote, should be avoided
- Functionality provided by the `state` command
- Subcommands available for listing resources, removing resources, moving items, pulling/pushing the state between local and remote backend, and rendering the current state



Listing Current Resources

“What resources are currently managed by Terraform?”

```
$ terraform state list
aws_instance.web_server
aws_ebs_volume.web_server_vol
```

See [command details](#)



Rendering Resource Attributes

“Show me the current configuration of this resource.”

```
$ terraform state show aws_instance.web_server
resource "aws_instance" "web_server" {
    ...
}
```

See [command details](#)



Renaming a Resource

“Renaming a resource without deleting/recreating it first”

```
$ terraform state mv aws_instance.web_server  
aws_instance.app_server  
Move "aws_instance.web_server" to "aws_instance.app_server"  
Successfully moved 1 object(s).
```

See [command details](#)



Remove an Item from the State

“I want to manage a resource outside of Terraform”

```
$ terraform state rm aws_instance.web_server
Removed aws_instance.web_server
Successfully removed 1 resource instance(s).
```

```
# Remove the resource from the configuration
$ terraform plan
$ terraform apply
```

See [command details](#)



EXERCISE

Interacting with the
State From the CLI



Enabling Verbose Logging

Rendering detailed information in error situations

- Terraform writes logs to standard output and offers the following log levels: TRACE, DEBUG, INFO, WARN, and ERROR
- The log level can be set via the environment variable `TF_LOG`
- The log level for Terraform itself and provider plugins can be controlled with the environment variable `TF_LOG_CORE` and `TF_LOG_PROVIDER`



Q & A



Understanding HCP Terraform Capabilities

How Does it Help in an Enterprise Setting?

What is HCP Terraform?

Commercial, cloud-hosted distribution of Terraform

Product	Hosting	Use Case
Terraform (Open Source)	Your machine/CI	Individual use, custom pipelines
HCP Terraform	HashiCorp SaaS	Teams who can use cloud-hosted
Terraform Enterprise	Self-hosted (your infrastructure)	Air-gapped, regulated industries, data sovereignty



HCP Terraform Features

Collaboration and governance capabilities

Capability	Description
Remote State	Secure, versioned state storage with locking
Remote Execution	Runs happen in consistent cloud environment
Team Collaboration	Shared workspaces with access controls
VCS Integration	Automatic runs triggered by Git commits
Policy as Code	Enforce rules with Sentinel or OPA
Private Registry	Share modules and providers internally



Open Source Terraform vs HCP

Open Source for individual use, HCP is ideal for teams

Aspect	Open-Source Terraform	HCP Terraform
State storage	Local or self-managed backend	Managed, encrypted, versioned
Execution	Your machine or CI/CD	HashiCorp's infrastructure
Collaboration	Manual coordination	Built-in workspaces and teams
Secrets	You manage	Encrypted variables
Audit trail	None	Full run history
Cost	Free	Free tier + paid plans



Feature: Account Management

Authentication and authorization

- A company can model organizational structure with the concepts *users, teams, and organizations*
- Permissions can be defined to control authorization of operations
- **Benefit:** Added security by only allowing authenticated access, User can only execute operations with granted permissions, Single Sign On (SSO) integration



Feature: VCS Integration

Terraform workflow integration as part of CI/CD automation

- Supports GitHub, GitLab, BitBucket, and Azure DevOps
- Automatic execution of a Terraform workflow upon commit
- Simplified release process of Terraform modules
- **Benefit:** Simplified setup of automatic execution for typical Terraform workflows against code changes without a lot of manual intervention



Feature: Workspaces

Not to be confused with local Terraform workspace

- HCP Terraform and local Terraform uses the same terminology, however, the meaning for a *workspace* is different
- In HCP Terraform a workspace is associated with a VCS repository
- Workspace settings include variable values, state, and secrets, and historical logs
- **Benefit:** Avoid local state file(s) and centralize state on remote



Feature: Config. Compliance

Enforced compliance as part of the plan-apply workflow

- Sentinel and Open Policy Agent (OPA) are compliance frameworks for centralizing company rules as code and enforcing them
- Triggers after terraform plan and before terraform apply commands and fails command if configuration doesn't follow policies
- **Benefit:** Standardizing on resource naming conventions or validating attribute values



Feature: Private Registry

Added security by only allowing members to access registry

- Enterprises can host their providers and modules in a private registry (only accessible to members of the organization)
- Includes functionality like versioning, and search capabilities like you already know from the [public Terraform registry](#)
- **Benefit:** Binaries cannot be modified by malicious attacker, attacker cannot publish a new version of binaries that contains malicious code



Other Notable Features

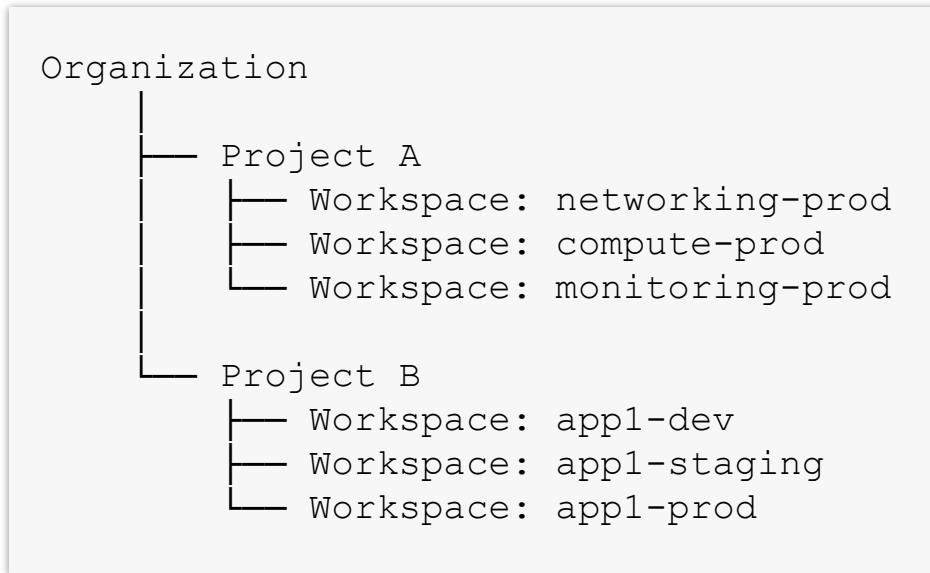
Refer to the [documentation](#) for more information

- API access: Ability to run Terraform operations via a RESTful API
- Remote operations: Consistent and reliable run environment with the help of virtual machines
- Cost estimation: Calculates expected cost when run on cloud provider like AWS, GCP, and Azure



HCP Terraform Hierarchy

Organization(s) > Project(s) > Workspace(s)



What are Organizations?

The root-level structure in HCP Terraform

- Most companies have ONE organization which maps to the company itself
- Some large enterprises have multiple (acquisitions, strict separation)
- Teams are defined at the org level, then granted access to projects/workspaces



Organization in UI

You'll be prompted to create a new organization to begin with

Organizations

+ Create organization

Terraform organizations let you manage organizations, projects, and teams.

Q Search by organization name

Organization name	Organization type	Actions
Terraform-Associate-Organization	 Terraform standalone	...

1–1 of 1

< 1 >

Items per page 10



What are Projects?

Projects group related workspaces together

- Projects are a relatively new feature (2023)
- Before projects, organizations had flat lists of workspaces
- Projects solve the "too many workspaces" problem
- Think of them as folders for workspaces



Projects in UI

To be created underneath an Organization

Terraform-Associate-Organization / Projects

Projects

+ New project

Projects in the Terraform-Associate-Organization organization. Create or select a project to get started.

Search by project name

Project name ↑↓	Description	Teams	Workspaces	Actions
Backend	No description	0	3	...
Default Project	No description	0	0	...

1–2 of 2

< 1 >

Items per page 20



Project Benefits

It's highly recommended to use projects

- **Focused View**
 - See only relevant workspaces
 - Filter by project in UI
- **Simplified Management**
 - Apply variable sets to all workspaces in a project
 - Assign team permissions at project level
 - Consistent policy enforcement



Project Benefits

It's highly recommended to use projects

- **Least Privilege**
 - Teams access only their projects
 - No organization-wide permissions needed



The Default Project

Every organization has a default project

- Cannot be deleted (can be renamed)
- New workspaces go here by default
- Catch-all for unassigned workspaces
- Best practice: Create specific projects and move workspaces out of “Default”



Default Project in UI

Automatically available with creation of an Organization

Terraform-Associate-Organization / Projects

Projects

+ New project

Projects in the Terraform-Associate-Organization organization. Create or select a project to get started.

Q Search by project name

Project name ↑↓	Description	Teams	Workspaces	Actions
<u>Default Project</u>	No description	0	0	...

1–1 of 1 < 1 > Items per page 20



Project Permissions

Teams can have different permissions per project

Permission	Capabilities	Use Case
Read	View workspaces and runs	View-only access, good for auditors
Plan	Queue plans	Can test changes but not apply - good for junior devs
Write	Queue applies	Full operational access
Admin	Manage workspace settings	Can change settings, manage variables



Project-Level Variable Sets

tfe provider interacts with resources supp. by HCP Terraform

```
resource "tfe_variable_set" "aws_creds" {
    name          = "AWS Credentials"
    organization = "my-org"
}

resource "tfe_project_variable_set" "platform_aws" {
    project_id      = tfe_project.platform.id
    variable_set_id = tfe_variable_set.aws_creds.id
}
```



What are Workspaces?

A *workspace* is an isolated environment containing:

- **Terraform configuration:** Your infrastructure code
- **State file:** Current state of managed resources
- **Variables:** Input values for the configuration
- **Run history:** Record of all plan/apply operations
- **Access controls:** Team permissions



Local Directories vs Workspaces

Key benefits: collaboration, audit trail, access control

Local Terraform	HCP Terraform
Separate directories	Separate workspaces
Local state files	Remote state storage
Manual collaboration	Built-in collaboration
No run history	Full audit trail
No access control	Full audit trail



Workspace Naming Conventions

Currently not enforced

Recommended format:

```
<application>-<layer>-<environment>
```

Examples:

- webapp-network-prod
- webapp-compute-staging
- api-database-dev



Workspaces in UI

Created with New button available to a Project

Terraform-Associate-Organization / Projects / Backend / Workspaces

Backend New

ID: prj-95cfmhGsg8mzhvswW

[Add project description](#)

Teams 0 Workspaces 3 Tags 0

Workspaces

Search by workspace name Needs attention Errorred Running On hold Completed All filters

No filters applied

Workspace name	Repository	Health	Latest change
api-database-dev No status reported	None	None	a few seconds ago
webapp-compute-staging No status reported	None	None	a few seconds ago
webapp-network-prod No status reported	None	None	a minute ago

1–3 of 3 Items per page 20



Q & A



EXERCISE

Apache Web Server
on AWS Accessible
via the Internet



Summary & Wrap Up

Last words of advice...



Thank you

